

Supplementary Materials to “A Generative Word Embedding Model and its Low Rank Positive Semidefinite Solution”

August 15, 2015

Proposition 1. *Suppose two words w_1, w_2 often collocate. In text snippets \dots, w_1, w_2, \dots where w_1 is in the focus and w_2 is in the context, and in \dots, w_2, w_1, \dots , where their roles exchange, there are probably common semantic regularities, by intuition. In the language of statistics, there should be positive correlation between $P(w_1|w_2)$ and $P(w_2|w_1)$. When using different representations for focus words and context words, this correlation is lost.*

Proof. Suppose focus words and context words use different representations \mathbf{u}_{w_i} and \mathbf{v}_{w_i} , and the two conditional probabilities are modeled as $P(w_1|w_2) = f(\mathbf{u}_{w_1}^\top \mathbf{v}_{w_2}, a_{w_1}, a_{w_2})$, and $P(w_2|w_1) = f(\mathbf{v}_{w_1}^\top \mathbf{u}_{w_2}, a_{w_2}, a_{w_1})$, where a_{w_i} is a non-interactive parameter. Then the respective contributions of the embeddings, $\mathbf{u}_{w_1}^\top \mathbf{v}_{w_2}$ and $\mathbf{v}_{w_1}^\top \mathbf{u}_{w_2}$, are completely irrelevant, other than being indirectly linked through the corpus statistics. \square

Fact 2. *Observations from the trained embeddings of various methods show that, the expectation of the embeddings is close to 0.*

Table 1 lists the expectations of the embeddings of popular methods. The expectation of the embeddings usually satisfies $\|E_{P(s)}[\mathbf{v}_s]\|_1 < \frac{1}{3}E_{P(s)}[\|\mathbf{v}_s\|_1]$, i.e. the magnitude of the expected embedding is much shorter than the expected embedding magnitude. This shows that the embeddings of different words point at “random” directions in the embedding space, and cancel each other to a large extent. This observation is natural since in the optimization objective, there is no explicit constraint on the directions of the embeddings.

Table 1: The expectations of the embeddings of word2vec, GloVe, Forest, PSD.

	word2vec	GloVe	Forest	PSD
$\ E_{P(s)}[\mathbf{v}_s]\ _1$	9.68	44.51	361.54	4.82
$E_{P(s)}[\ \mathbf{v}_s\ _1]$	31.38	113.46	1711.4	39.02

In “word2vec” [2], there is a normalizing constant for each focus word s_j in the embedding function. This is equivalent to adopting a constant residual r_{s_j} . In GloVe [3], the residual of s_i, s_j is a linear combination of the residuals of words s_i and s_j . However we will prove that a constant residual, or even the linear combination of the residuals of s_i and s_j , could not satisfy Bayes’s theorem, given the assumption that $E_{P(s)}[\mathbf{v}_s] \approx 0$. Therefore we have to generalize the residual to a residual for each bigram s_i, s_j , denoted as $a_{s_i s_j}$.

References

- [1] S. LI, J. ZHU, AND C. MIAO, *A generative word embedding model and its low rank positive semidefinite solution*, in Proceedings of EMNLP 2015, 2015.
- [2] T. MIKOLOV, I. SUTSKEVER, K. CHEN, G. S. CORRADO, AND J. DEAN, *Distributed representations of words and phrases and their compositionality*, in Proceedings of NIPS 2013, 2013, pp. 3111–3119.
- [3] J. PENNINGTON, R. SOCHER, AND C. D. MANNING, *Glove: Global vectors for word representation*, Proceedings of the Empirical Methods in Natural Language Processing (EMNLP 2014), 12 (2014).
- [4] N. TIMME, W. ALFORD, B. FLECKER, AND J. M. BEGGS, *Synergy, redundancy, and multivariate information measures: an experimentalist’s perspective*, Journal of computational neuroscience, 36 (2014), pp. 119–140.
- [5] P. L. WILLIAMS AND R. D. BEER, *Nonnegative decomposition of multivariate information*, CoRR, abs/1004.2515 (2010).
- [6] D. YOGATAMA, M. FARUQUI, C. DYER, AND N. A. SMITH, *Learning word representations with hierarchical sparse coding*, in Proceedings of ICML, 2015.